

# An empirical study of the naïve REINFORCE algorithm for predictive maintenance of industrial machines

Rajesh Siraskar

June 14, 2023

## Abstract

”Occam’s razor is a principle often attributed to 14th century friar William of Ockham that says that if you have two competing ideas to explain the same phenomenon, you should prefer the simpler one.” - Chris Simms (NewScientist, )

Everything should be made as simple as possible, but not simpler. – Albert Enstien

Our broad goal is to understand the performance of this naïve algorithm with pitched against the more advanced algorithms, specifically the DQN, A2C and PPO.

Our approach uses simulated as well as real data. Simple as well as complex env. Without as well as with noise. Our systematic analysis and study show that the naïve REINFORCE implementation out performs these algorithms by xx, xx, and xx respectively. The variance as measured by std. dev is xx, xx, xx

## 1 Introduction

Introduced in 1992, the REINFORCE algorithm is considered as a basic reinforcement learning algorithm. It is a policy-based, on-policy as well as off-policy algorithm, capable of handling both discrete and continuous observation and action domains.

In practice the REINFORCE algorithm is considered as “weak” learner and superseded by several algorithms developed since. Most notably the Q-Learning

and its deep-neural network version, the DQN, followed by Actor-Critic and one of the most robust modern day algorithms, the PPO.

[Reinforcement Learning in Robotics: A Survey](#) - Jens Kober J. Andrew Bagnell Jan Peters - Initial gradient-based approaches such as finite differences gradients or REINFORCE (Williams, 1992) have been rather slow. The weight perturbation algorithm is related to REINFORCE but can deal with non-Gaussian distributions which significantly improves the signal to noise ratio of the gradient (Roberts et al., 2010). Recent natural policy gradient approaches (Peters and Schaal, 2008c,b) have allowed for faster convergence which may be advantageous for robotics as it reduces the learning time and required real-world interactions.

xxx  
xxx  
xxx  
xxx  
xxx

## 1.1 Algorithm timelines

- 1947: Monte Carlo Sampling
- 1959: Temporal Difference Learning
- 1989: Q-Learning
- 1992: REINFORCE
- 2013: DQN
- 2016: A3C
- 2017: PPO

## 2 The REINFORCE algorithm

Three key features of any RL algorithm:

1. Policy:  $\pi_\theta$  = Probabilities of all actions, given a state. Parameterized by  $\theta$
2. Objective function:

$$\max_{\theta} J(\pi_\theta) = \mathbb{E}_{\tau \sim \pi_\theta} [R(\tau)] \quad (1)$$

3. Method: Way to update the parameters = Policy Gradient

## 2.1 Policy gradient numerical computation

1. Plain vanilla:

$$\nabla_{\theta} J(\pi_{\theta}) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^T R_t(\tau) \nabla_{\theta} \ln \pi_{\theta}(a_t | s_t) \right] \quad (2)$$

2. With Monte Carlo sampling and approximation:  $\nabla_{\theta} J(\pi_{\theta}) \approx \left[ \sum_{t=0}^T R_t(\tau) \nabla_{\theta} \ln \pi_{\theta}(a_t | s_t) \right]$
3. With baseline:  $\nabla_{\theta} J(\theta) \approx \left[ \sum_{t=0}^T (R_t(\tau) - b(s_t)) \nabla_{\theta} \ln \pi_{\theta}(a_t | s_t) \right]$
4. Where, baseline does not change per time-step, it is for the entire trajectory
5. One baseline option:  $V^{\pi}$  - leads to Actor-Critic algorithm
6. Simpler option: Average returns over trajectory:  $b = \frac{1}{T} \sum_{t=0}^T R_t(\tau)$

Algorithm

## 3 About Stable-Baselines-3

- SB3- paper ([Raffin et al., 2021](#)), [Raffin et al. \(2021\)](#)
- sb-3 main doc – ([SB3, b](#))
- sb-3 ppo doc – ([SB3, a](#))

## 4 Method

We normalize the tool wear and other state features,  $x \in [0, 1] \subset \mathbb{R}$ . This allows for adding white noise of similar magnitudes across experiments of different data-sets

## 5 Results

## 6 Discussion

## 7 Conclusion

## References

Stable-baselines3 docs - reliable reinforcement learning implementations, a.

URL <https://stable-baselines3.readthedocs.io/en/master/modules/ppo.html#how-to-replicate-the-results>.

Ppo, b. URL <https://stable-baselines3.readthedocs.io/en/master/index.html>. Accessed: 2023-05-14.

Antonin Raffin, Ashley Hill, Adam Gleave, Anssi Kanervisto, Maximilian Ernestus, and Noah Dormann. Stable-baselines3: Reliable reinforcement learning implementations. *J. Mach. Learn. Res.*, 22(1), jan 2021. ISSN 1532-4435.